

Read-only diskless filesystem
with Debian 7.0
CTRL.01.01.001

Sébastien BLANCHET

Oct 13, 2014

Contents

1	Build NFS root	5
1.1	Create NFS root	5
1.1.1	Create the host computer	5
1.1.2	Create a bootstrap	5
1.1.3	/etc/fstab	5
1.1.4	Directory /etc/default	6
1.1.5	/etc/apt/sources.list	6
1.1.6	/etc/hosts	6
1.1.7	/bin/whereami	6
1.1.8	Configure /etc/mtab	7
1.1.9	Configure reboot	7
1.1.10	Configure the root user	7
1.1.11	Install additionnal firmware	7
1.1.12	Install additionnal driver	8
1.1.13	Build a PXE initrd	8
1.1.14	Driver blacklisting	8
1.2	Prepare NFS and TFTP server	9

Chapter 1

Build NFS root

1.1 Create NFS root

1.1.1 Create the host computer

A Debian host computer is required to create the nfsroot filesystem. In theory, you can use any Debian computer, whatever the release or the architecture are. But in practice it is more convenient to use the same release and architecture for both. So create a new virtual machine with VirtualBox and install Debian-7.0-amd64 into.

1.1.2 Create a bootstrap

debootstrap creates a basic Debian system.

```
# aptitude -y install debootstrap
# mkdir /home/nfsroot
# debootstrap wheezy /home/nfsroot http://ftp2.fr.debian.org/debian
```

The mirror URL is optional, but it downloads faster if you select a mirror near you. See manual `debootstrap(8)` for details.

Now we customize this basic Debian system in `/home/nfsroot`.

```
# export NFSROOT=/home/nfsroot
```

1.1.3 `/etc/fstab`

Listing 1.1: `$NFSROOT/etc/fstab`

proc	/proc	proc	defaults	0 0
/dev/nfs	/	nfs	tcp,nolock	0 0
none	/tmp	tmpfs	defaults	0 0

none	/var/tmp	tmpfs	defaults	0 0
none	/media	tmpfs	defaults	0 0
none	/var/log	tmpfs	defaults	0 0
nfssrv:/path/to/home	/home	nfs	tcp,nolock	1 2

The `nolock` `nfs` option is very important, otherwise `sqlite` does not work on these mount points. See `man 5 nfs`, for other `nfs` options.

1.1.4 Directory `/etc/default`

See manuals `rcS(5)` and `tmpfs(5)` for details.

```
# echo ASYNCMOUNTNFS=no >> ${NFSROOT}/etc/default/rcS
# echo RAMTMP=yes >> ${NFSROOT}/etc/default/tmpfs
```

1.1.5 `/etc/apt/sources.list`

Point to the nearest mirror, and add `contrib` and `non-free` repositories.

Listing 1.2: `/etc/apt/sources.list`

```
deb http://ftp2.fr.debian.org/debian/ wheezy main contrib non-free
deb-src http://ftp2.fr.debian.org/debian/ wheezy main contrib non-free

deb http://security.debian.org/ wheezy/updates main contrib non-free
deb-src http://security.debian.org/ wheezy/updates main contrib non-free
```

Host and `nfsroot` must share the same configuration:

```
# cp /etc/apt/sources.list ${NFSROOT}/etc/apt/sources.list
```

1.1.6 `/etc/hosts`

Remove all IPv6 lines and add IP for each NFS client.

Listing 1.3: `${NFSROOT}/etc/hosts`

127.0.0.1	localhost
192.168.0.186	pc-client1
192.168.0.187	pc-client2
192.168.0.188	pc-client3

1.1.7 `/bin/whereami`

Create a script `/bin/whereami` to setup client hostname from `/etc/hosts`

Listing 1.4: \$NFSROOT/bin/whereami

```
#!/bin/bash
#finds node's hostname based on matching ip in /etc/hosts

grep 'ifconfig' \
    | grep 'inet addr:' \
    | grep -v '127.0.0.1' \
    | /usr/bin/cut -d: -f2 \
    | /usr/bin/awk '{ print $1}'' /etc/hosts \
    | /usr/bin/awk '{print $2}'
```

Add executable bit:

```
# chmod +x ${NFSROOT}/bin/whereami
```

Modify \$NFSROOT/etc/init.d/hostname.sh to replace

```
[ -f /etc/hostname ] && HOSTNAME="$(cat /etc/hostname)"
```

by

```
[ -f /etc/hostname ] && HOSTNAME="$(/bin/whereami)"
```

1.1.8 Configure /etc/mtab

```
# ln -s /proc/mounts ${NFSROOT}/etc/mtab
```

1.1.9 Configure reboot

Tweak reboot scripts to avoid cutting down the network interface.

```
# sed -i.orig -e 's/reboot -d ./reboot -d -f/g' \
    ${NFSROOT}/etc/init.d/reboot
# sed -i.orig -e 's/halt -d ./reboot -d -f/g' \
    ${NFSROOT}/etc/init.d/halt
```

1.1.10 Configure the root user

```
# chroot /home/introot
# passwd
```

1.1.11 Install additionnal firmware

It must be done before rebuilding initrd.

```
# aptitude update
# aptitude install firmware-linux
```

1.1.12 Install additionnal driver

It must be done before rebuilding initrd.

In this example we add an updated driver to support Realtek 8111G.

Download `r8168-8.037.00.tar.bz2` from <http://www.realtek.com>

```
# tar xf r8168-8.037.00.tar.bz2
# cd r8168-8.037.00
# bash ./autorun.sh
```

The `autorun` script does:

- build the new driver
- disable the linux driver `r8169`
- install the new driver `r8168` in `/lib/modules/`uname -r`` so it will be automatically present when we rebuild the `initrd`.

1.1.13 Build a PXE initrd

The default `initrd` image is not suitable for diskless computer, so rebuild a new one.

```
# mkdir -p /var/lib/tftpboot/
# cp -r /etc/initramfs-tools /etc/initramfs-pxe
# echo BOOT=nfs >> /etc/initramfs-pxe/initramfs.conf
# mkinitramfs -d /etc/initramfs-pxe \
    -o /var/lib/tftpboot/initrd.pxe-`uname -r` `uname -r`
```

Initrd customization

You can uncompress the `initrd` image to edit its content:

```
# mkdir initrd
# cd initrd
# gunzip2 -dc initrd.pxe-`uname -r` | cpio -i
# #MODIFY the files here
# find . | cpio -c -o | gzip -9 > initrd.pxe-`uname -r`
```

1.1.14 Driver blacklisting

During a Linux diskless booting process, the NFS root filesystem **must** be mounted through the PXE booting interface **only**. But during the `initrd` step, the kernel does not know which interface has been used for PXE booting, so it will use the first detected interface instead.

For example if you add a Broadcom network card into a Beckhoff computer, you may encounter this kind of problem:

- the computer can boot only from its internal Intel network interfaces.
- `initrd` detects the Broadcom before the Intel network interfaces, so it tries to mount the NFS root filesystem through the wrong interface. So it fails.

The solution is to blacklist the broadcom driver `tg3` during the init. There are two ways to do it:

- by adding `blacklist=tg3` to the kernel command line in the `pxeconfig` file. It is the easier and more flexible solution.
- by editing the `blacklist.conf` file inside the `initrd`

Editing `blacklist.conf` inside `initrd` file

```
# mkdir initrd
# cd initrd
# bunzip2 -dc initrd.pxe-'uname -r' | cpio -i
# echo "blacklist tg3" >> etc/modprobe.d/blacklist.conf
# find . | cpio -c -o | bzip2 -9 > initrd.pxe-'uname -r'
```

Note: With the Debian 7.0 i386 3.2.0-4-rt-686-pae kernel, this procedure does not work. The kernel is unable to mount the new `initrd` (it works only with the `initrd` created with `mkinitramfs`).

1.2 Prepare NFS and TFTP server

- Transfer `nfsroot` to the `nfs` server.
- Transfer kernel and `initrd` to the `tftp` server.

Create a `pxeconfig` file

Listing 1.5: `$NFSROOT/etc/hosts`

```
# boot diskless computer with debian wheezy
default menu.c32
prompt 0
menu title pc-client

ontimeout linux-3.2.0-3-amd64-wheezy
timeout 50

label linux-3.2.0-3-amd64-wheezy
  menu label linux-3.2.0-3-amd64 wheezy
  kernel vmlinuz-3.2.0-3-amd64
```

```
append root=/dev/nfs initrd=initrd.pxe-3.2.0-3-amd64 \
        nfsroot=nfsserver:/path/to/nfsroot ro panic=60
ipappend 1
```

Note: the backslash in the **append** line is only for printing purpose, to show that the line continue. But in reality, all the arguments *must* be on the same line.